

ENV 797 - Time Series Analysis for Energy and Environment

Applications | Spring 2026

Assignment 5 - Due date 02/17/26

Matthew Fischer

Directions

You should open the .rmd file corresponding to this assignment on RStudio. The file is available on our class repository on Github.

Once you have the file open on your local machine the first thing you will do is rename the file such that it includes your first and last name (e.g., “LuanaLima_TSA_A05_Sp26.Rmd”). Then change “Student Name” on line 3 with your name.

Then you will start working through the assignment by **creating code and output** that answer each question. Be sure to use this assignment document. Your report should contain the answer to each question and any plots/tables you obtained (when applicable).

When you have completed the assignment, **Knit** the text and code into a single PDF file. Submit this pdf using Canvas.

R packages needed for this assignment: “readxl”, “ggplot2”, “forecast”, “tseries”, and “Kendall”. Install these packages, if you haven’t done yet. Do not forget to load them before running your script, since they are NOT default packages.\

```
#Load/install required package here
library(forecast)

## Registered S3 method overwritten by 'quantmod':
##   method           from
##   as.zoo.data.frame zoo

library(tseries)
library(ggplot2)
library(Kendall)
library(lubridate)

##
## Attaching package: 'lubridate'
## The following objects are masked from 'package:base':
## 
##   date, intersect, setdiff, union
library(tidyverse) #load this package so you can clean the data frame using pipes

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr    1.1.4    v stringr 1.6.0
## vforcats 1.0.1    v tibble   3.3.1
## v purrr   1.2.1    v tidyrr   1.3.2
```

```

## v readr  2.1.6
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
library(readxl)

```

Consider the same data you used for A04 from the spreadsheet “Table_10.1_Renewable_Energy_Production_and_Consumption_by_Source.xlsx”. The data comes from the US Energy Information and Administration and corresponds to the December 2025 Monthly Energy Review.

```

#Importing data set - using readxl package
energy_data <- read_excel(
  path=".~/Data/Table_10.1_Renewable_Energy_Production_and_Consumption_by_Source.xlsx",
  skip = 12,
  sheet="Monthly Data",
  col_names=FALSE
)

## New names:
## * `` -> `...1`
## * `` -> `...2`
## * `` -> `...3`
## * `` -> `...4`
## * `` -> `...5`
## * `` -> `...6`
## * `` -> `...7`
## * `` -> `...8`
## * `` -> `...9`
## * `` -> `...10`
## * `` -> `...11`
## * `` -> `...12`
## * `` -> `...13`
## * `` -> `...14`

#Now let's extract the column names from row 11 only
read_col_names <- read_excel(
  path=".~/Data/Table_10.1_Renewable_Energy_Production_and_Consumption_by_Source.xlsx",
  skip = 10,
  n_max = 1,
  sheet="Monthly Data",
  col_names=FALSE
)

## New names:
## * `` -> `...1`
## * `` -> `...2`
## * `` -> `...3`
## * `` -> `...4`
## * `` -> `...5`
## * `` -> `...6`
## * `` -> `...7`
## * `` -> `...8`
## * `` -> `...9`
## * `` -> `...10`

```

```

## * `` -> `...11`
## * `` -> `...12`
## * `` -> `...13`
## * `` -> `...14`

colnames(energy_data) <- read_col_names
nobs <- nrow(energy_data)

nobs=nrow(energy_data)
nvar=ncol(energy_data)

head(energy_data)

## # A tibble: 6 x 14
##   Month           `Wood Energy Production` `Biofuels Production` 
##   <dttm>          <dbl> <chr>
## 1 1973-01-01 00:00:00    130. Not Available
## 2 1973-02-01 00:00:00    117. Not Available
## 3 1973-03-01 00:00:00    130. Not Available
## 4 1973-04-01 00:00:00    125. Not Available
## 5 1973-05-01 00:00:00    130. Not Available
## 6 1973-06-01 00:00:00    125. Not Available
## # i 11 more variables: `Total Biomass Energy Production` <dbl>,
## #   `Total Renewable Energy Production` <dbl>,
## #   `Hydroelectric Power Consumption` <dbl>,
## #   `Geothermal Energy Consumption` <dbl>, `Solar Energy Consumption` <chr>,
## #   `Wind Energy Consumption` <chr>, `Wood Energy Consumption` <dbl>,
## #   `Waste Energy Consumption` <dbl>, `Biofuels Consumption` <chr>,
## #   `Total Biomass Energy Consumption` <dbl>, ...

```

Handling Missing Data

Q1

Using the original dataset, create a new data frame that includes only the following variables: **Date**, **Solar Energy Consumption** and **Wind Energy Consumption**. Check the class of columns, you will see that they are stored are characters instead of numbers. Because solar generation begins later in the sample, the early observations are recorded as “Not Available”. Convert the data to numeric, the “Not Available” will became NAs.

You may either filter out the “Not Available” rows and then convert the column to numeric or convert first and then remove missing values using drop_na() (or na.omit()). If you are comfortable using pipes for data wrangling, please do so.

Important: Note that we dropping the missing observations instead of interpolating is because they only happen in the beginning of the series!

```

energy_subset <- energy_data |>
  select(
    Date = Month,
    Solar_Energy_Consumption = `Solar Energy Consumption`,
    Wind_Energy_Consumption = `Wind Energy Consumption`)
  ) |>
  mutate(
    Solar_Energy_Consumption = as.numeric(Solar_Energy_Consumption),
    Wind_Energy_Consumption = as.numeric(Wind_Energy_Consumption))

```

```
) |>
drop_na()

## Warning: There were 2 warnings in `mutate()` .
## The first warning was:
## i In argument: `Solar_Energy_Consumption =
##   as.numeric(Solar_Energy_Consumption)` .
## Caused by warning:
## ! NAs introduced by coercion
## i Run `dplyr::last_dplyr_warnings()` to see the 1 remaining warning.

str(energy_subset)

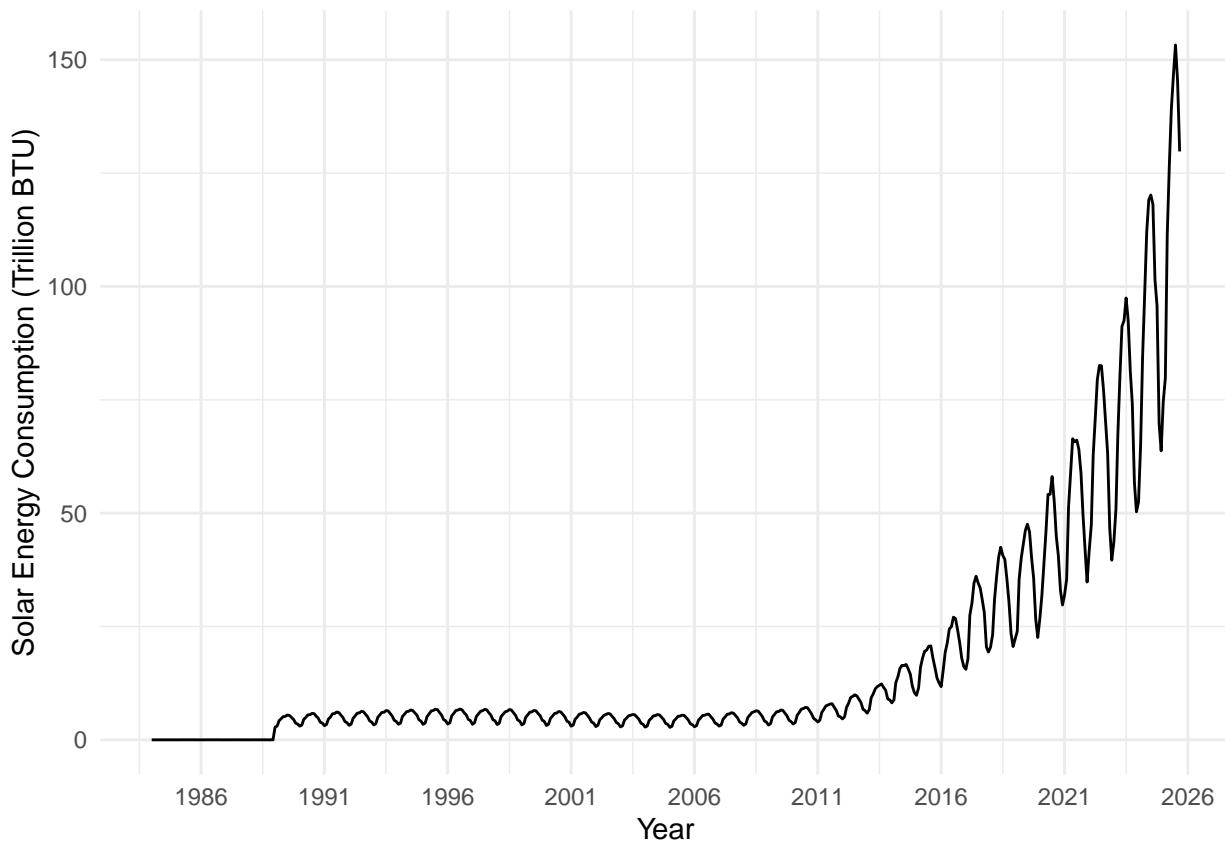
## tibble [501 x 3] (S3:tbl_df/tbl/data.frame)
## $ Date : POSIXct[1:501], format: "1984-01-01" "1984-02-01" ...
## $ Solar_Energy_Consumption: num [1:501] 0 0 0.001 0.001 0.002 0.003 0.001 0.003 0.003 0.002 ...
## $ Wind_Energy_Consumption : num [1:501] 0 0.001 0.001 0.002 0.003 0.002 0.002 0.001 0.002 0.003 ...
```

Q2

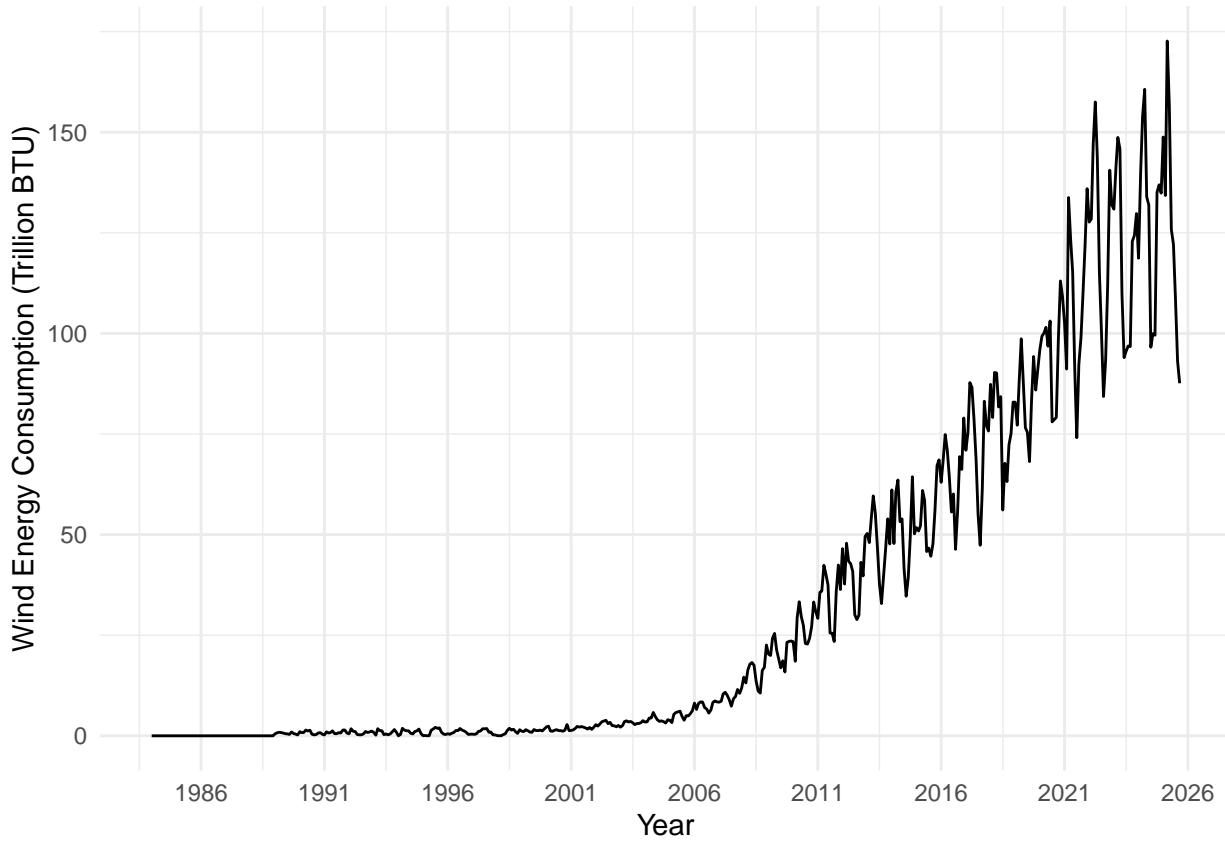
Plot the Solar and Wind energy consumption over time using ggplot. Plot each series on a separate graph. No need to add legend. Add informative names to the y axis using `ylab()`. Explore the function `scale_x_date()` on ggplot and see if you can change the x axis to improve your plot. Hint: use `scale_x_date(date_breaks = "5 years", date_labels = "%Y")`)

```
energy_subset$Date <- as.Date(energy_subset$Date)
```

```
ggplot(energy_subset, aes(x = Date, y = Solar_Energy_Consumption)) +  
  geom_line() +  
  ylab("Solar Energy Consumption (Trillion BTU)") +  
  xlab("Year") +  
  scale_x_date(date_breaks = "5 years", date_labels = "%Y") +  
  theme_minimal()
```



```
ggplot(energy_subset, aes(x = Date, y = Wind_Energy_Consumption)) +  
  geom_line() +  
  ylab("Wind Energy Consumption (Trillion BTU)") +  
  xlab("Year") +  
  scale_x_date(date_breaks = "5 years", date_labels = "%Y") +  
  theme_minimal()
```



Q3

Now plot both series in the same graph, also using ggplot(). Use function `scale_color_manual()` to manually add a legend to ggplot. Make the solar energy consumption red and wind energy consumption blue. Add informative name to the y axis using `ylab("Energy Consumption")`. And use function `scale_x_date()` to set x axis breaks every 5 years.

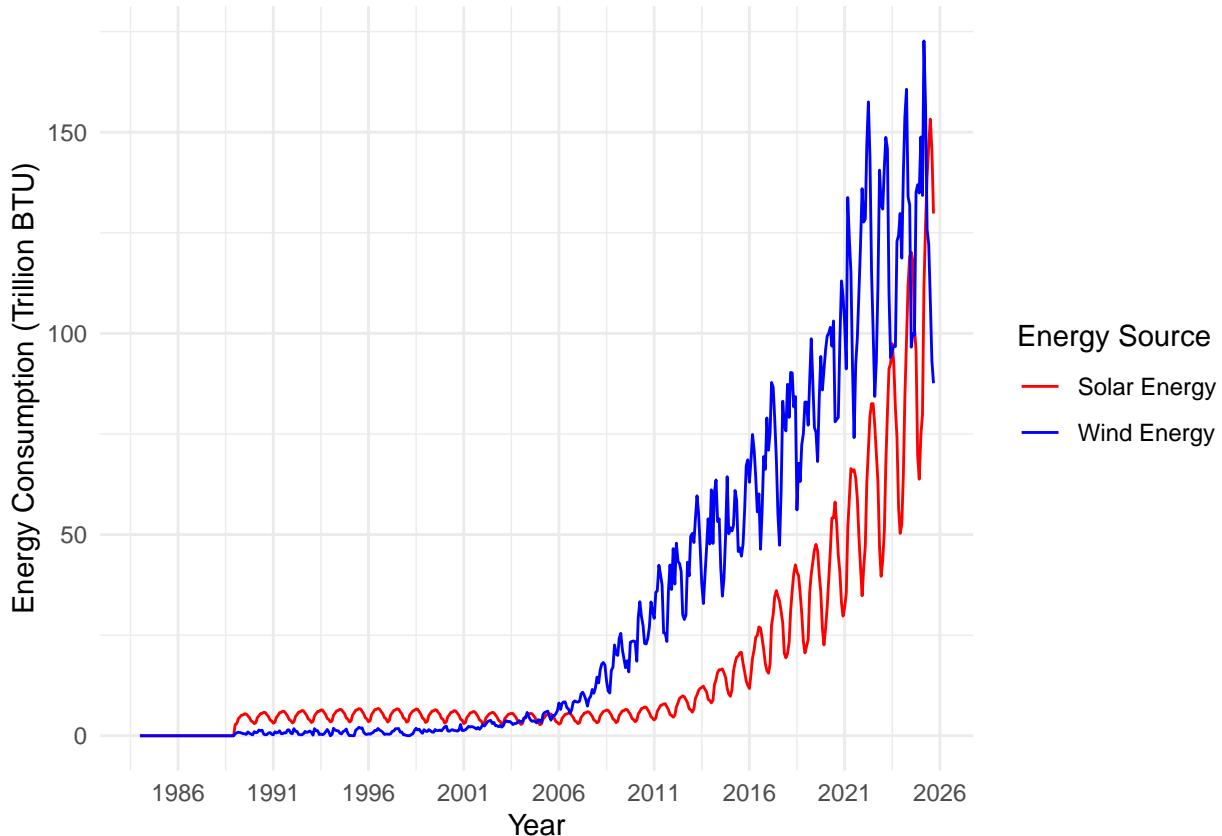
```
energy_long <- energy_subset %>%
  pivot_longer(
    cols = c(Solar_Energy_Consumption, Wind_Energy_Consumption),
    names_to = "Energy_Type",
    values_to = "Consumption"
  )

ggplot(energy_long, aes(x = Date, y = Consumption, color = Energy_Type)) +
  geom_line() +
  ylab("Energy Consumption (Trillion BTU)") +
  xlab("Year") +
  scale_color_manual(
    values = c(
      "Solar_Energy_Consumption" = "red",
      "Wind_Energy_Consumption" = "blue"
    ),
    labels = c(
      "Solar Energy",
      "Wind Energy"
    )
  ),
```

```

    name = "Energy Source"
) +
scale_x_date(date_breaks = "5 years", date_labels = "%Y") +
theme_minimal()

```



Decomposing the time series

The stats package has a function called `decompose()`. This function only take time series object. As the name says the decompose function will decompose your time series into three components: trend, seasonal and random. This is similar to what we did in the previous script, but in a more automated way. The random component is the time series without seasonal and trend component.

Additional info on `decompose()`.

- 1) You have two options: alternative and multiplicative. Multiplicative models exhibit a change in frequency over time.
- 2) The trend is not a straight line because it uses a moving average method to detect trend.
- 3) The seasonal component of the time series is found by subtracting the trend component from the original data then grouping the results by month and averaging them.
- 4) The random component, also referred to as the noise component, is composed of all the leftover signal which is not explained by the combination of the trend and seasonal component.

Q4

Transform wind and solar series into a time series object and apply the decompose function on them using the additive option, i.e., `decompose(ts_data, type = "additive")`. What can you say about the trend component? What about the random component? Does the random component look random? Or does it appear to still have some seasonality on it?

```

solar_ts <- ts(
  energy_subset$Solar_Energy_Consumption,
  start = c(1984, 1),
  frequency = 12
)

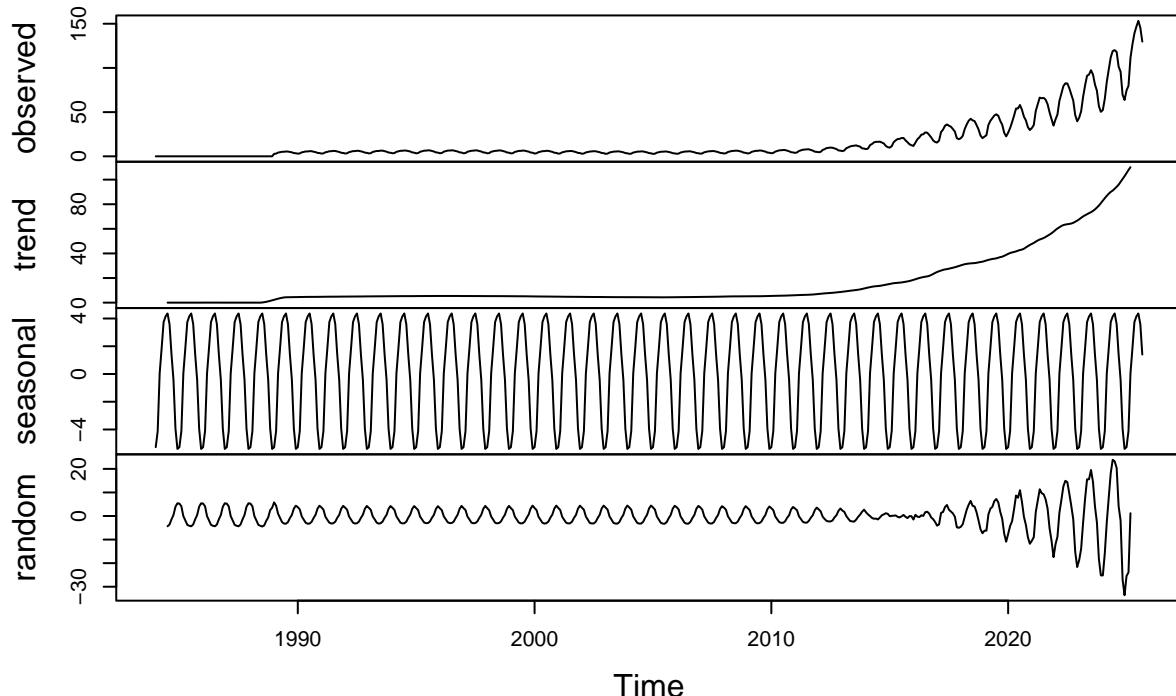
wind_ts <- ts(
  energy_subset$Wind_Energy_Consumption,
  start = c(1984, 1),
  frequency = 12
)

solar_add_decomp <- decompose(solar_ts, type = "additive")
wind_add_decomp <- decompose(wind_ts, type = "additive")

plot(solar_add_decomp)

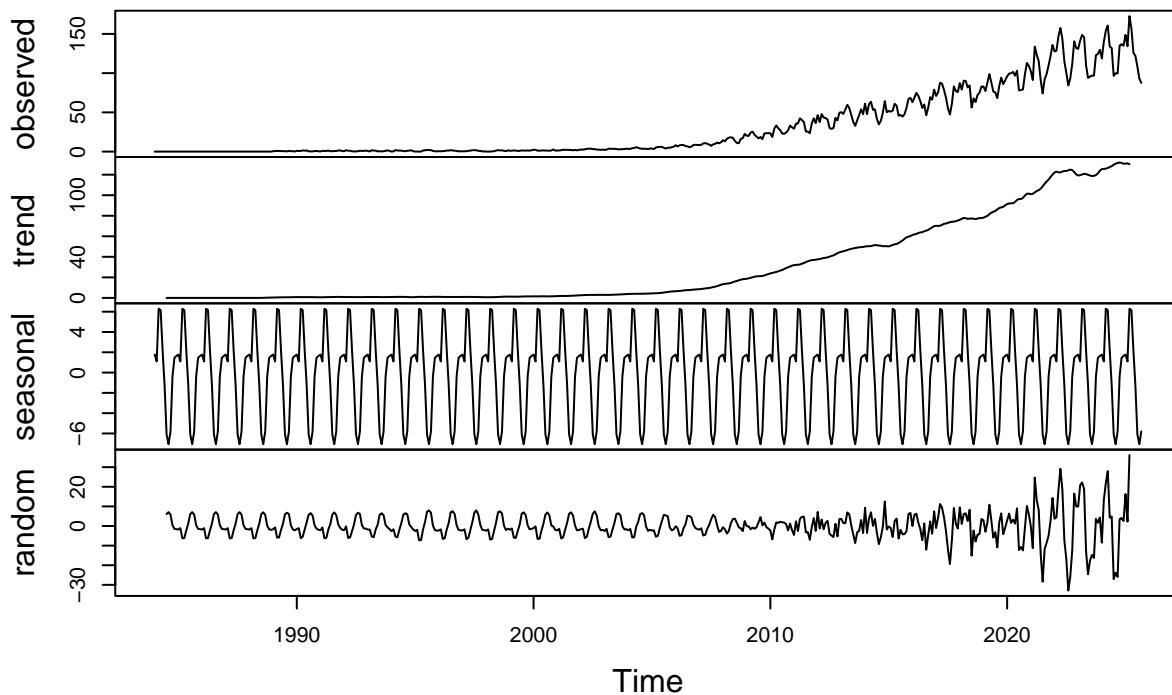
```

Decomposition of additive time series



```
plot(wind_add_decomp)
```

Decomposition of additive time series



Answer: The trends for both look about the same, increasing exponentially, although the wind energy consumption seems to be evening out around 2020. Both random components seem to still have seasonality.

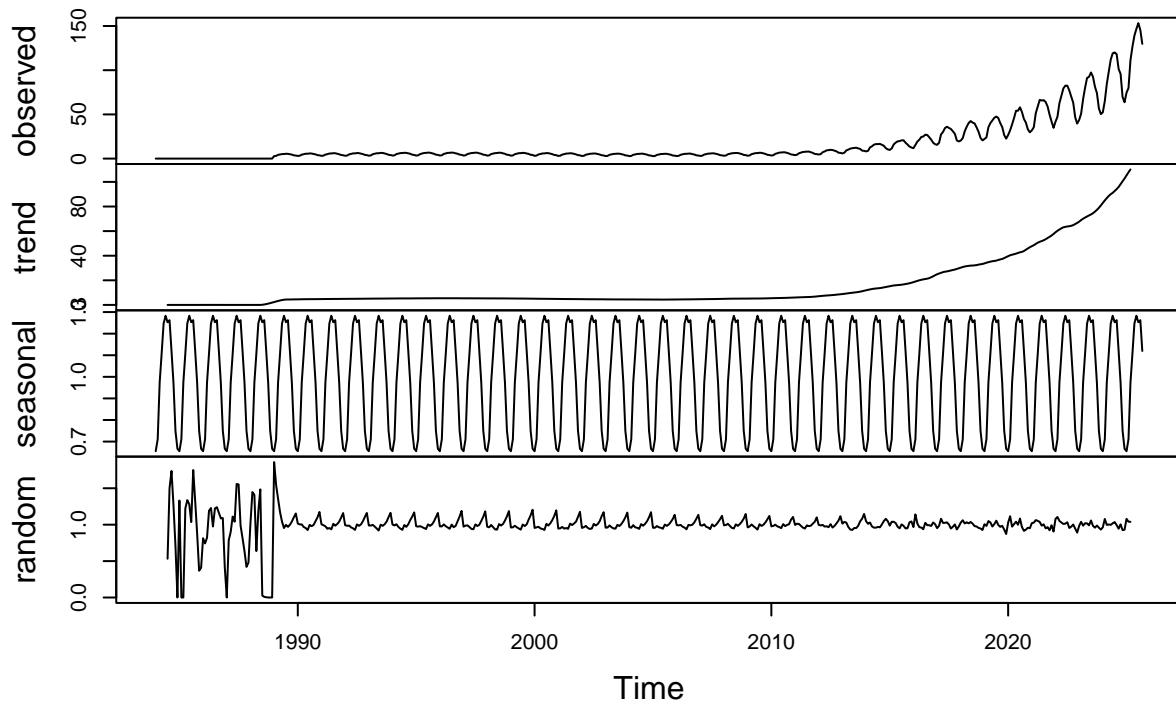
Q5

Use the decompose function again but now change the type of the seasonal component from additive to multiplicative. What happened to the random component this time?

```
solar_mult_decomp <- decompose(solar_ts, type = "multiplicative")
wind_mult_decomp <- decompose(wind_ts, type = "multiplicative")

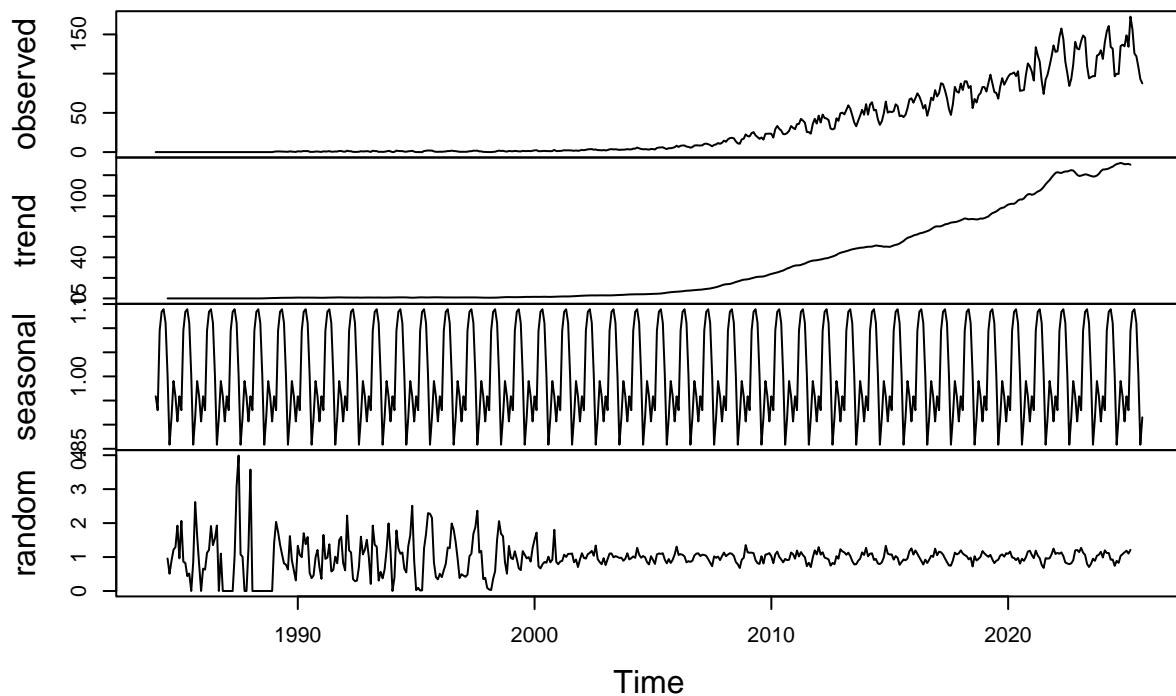
plot(solar_mult_decomp)
```

Decomposition of multiplicative time series



```
plot(wind_mult_decomp)
```

Decomposition of multiplicative time series



Answer: The random components look more random now, although there still seems to be seasonality from around 1990 to present for solar energy consumption and from around 2000 to present for wind energy consumption.

Q6

When fitting a model to this data, do you think you need all the historical data? Think about the data from 80s, 90s and early 20s. Are there any information from those years we might need to forecast the next six months of Solar and/or Wind consumption. Explain your response.

Answer: We probably do not need all the historical data, the overall pattern now seems to be different than the early 2000s and before, and the seasonal pattern seems to be stable enough that we do not need the historical data to identify it.

Q7

Create a new time series object where historical data starts on January 2014. Hint: use `filter()` function so that you don't need to point to row numbers, i.e, `filter(xxxx, year(Date) >= 2014)`. Apply the decompose function `type=additive` to this new time series. Comment on the results. Does the random component look random?

```
energy_recent <- energy_subset |>
  filter(year(Date) >= 2014)

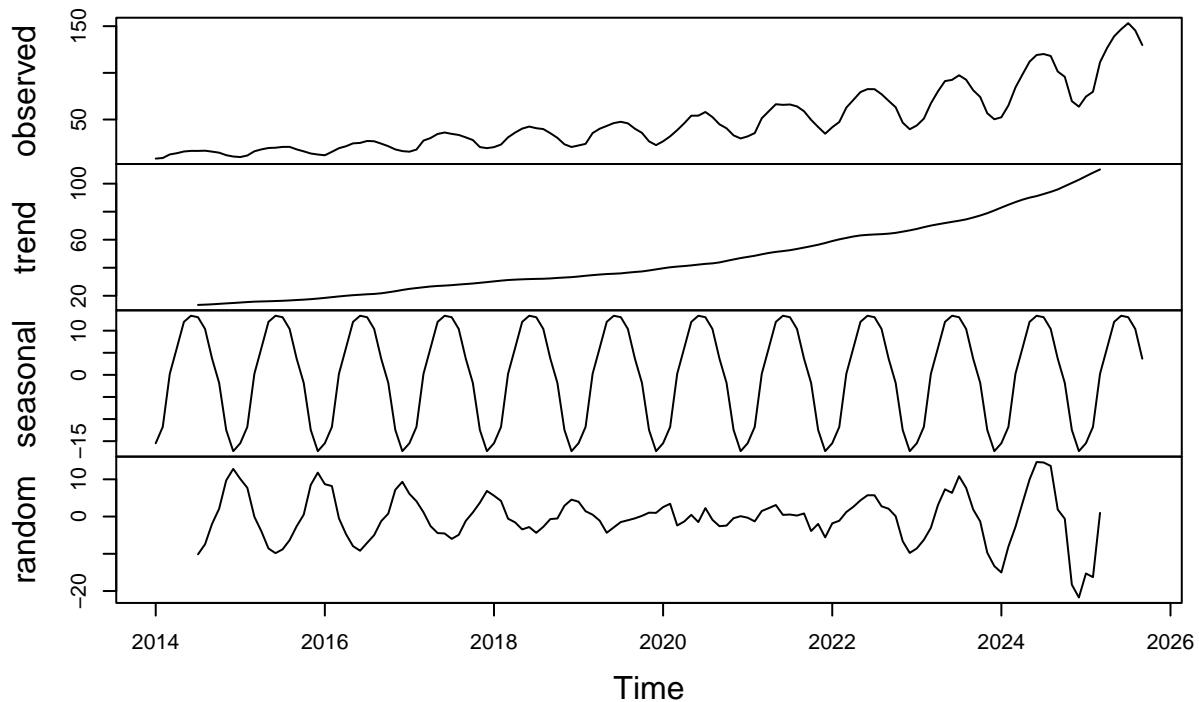
solar_recent_ts <- ts(
  energy_recent$Solar_Energy_Consumption,
  start = c(2014, 1),
  frequency = 12
)

wind_recent_ts <- ts(
  energy_recent$Wind_Energy_Consumption,
  start = c(2014, 1),
  frequency = 12
)

solar_recent_decomp <- decompose(solar_recent_ts, type = "additive")
wind_recent_decomp <- decompose(wind_recent_ts, type = "additive")

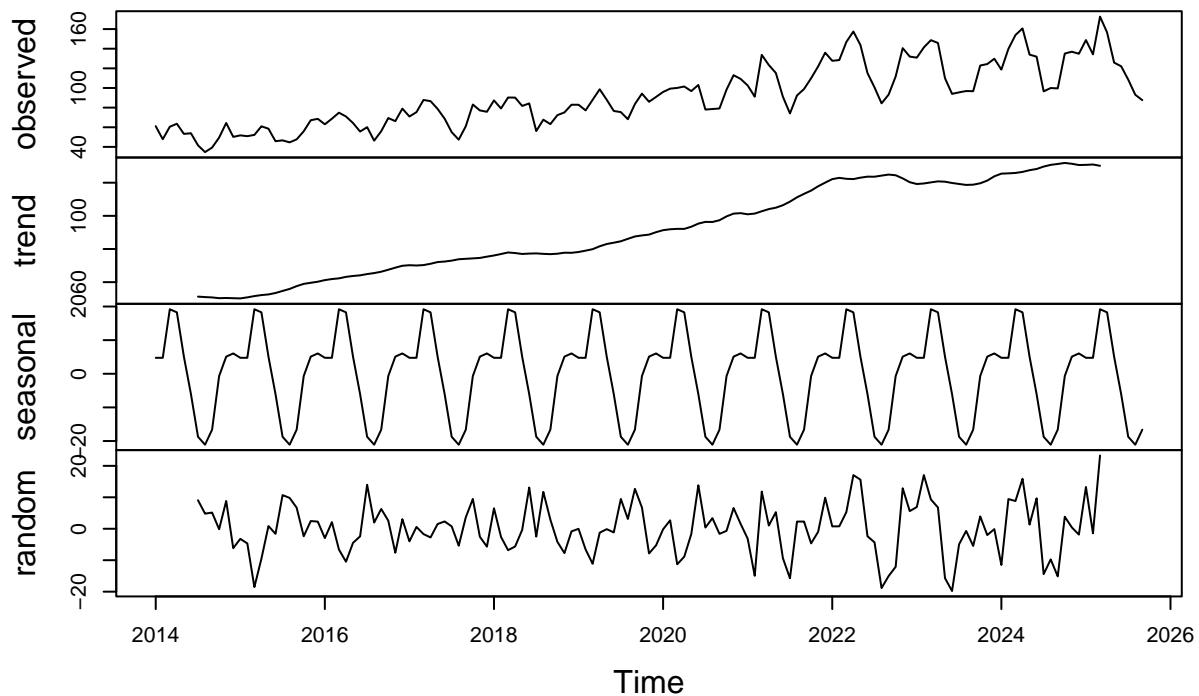
plot(solar_recent_decomp)
```

Decomposition of additive time series



```
plot(wind_recent_decomp)
```

Decomposition of additive time series



Answer: The random components seem more random than before but there still seems to be some seasonality present.

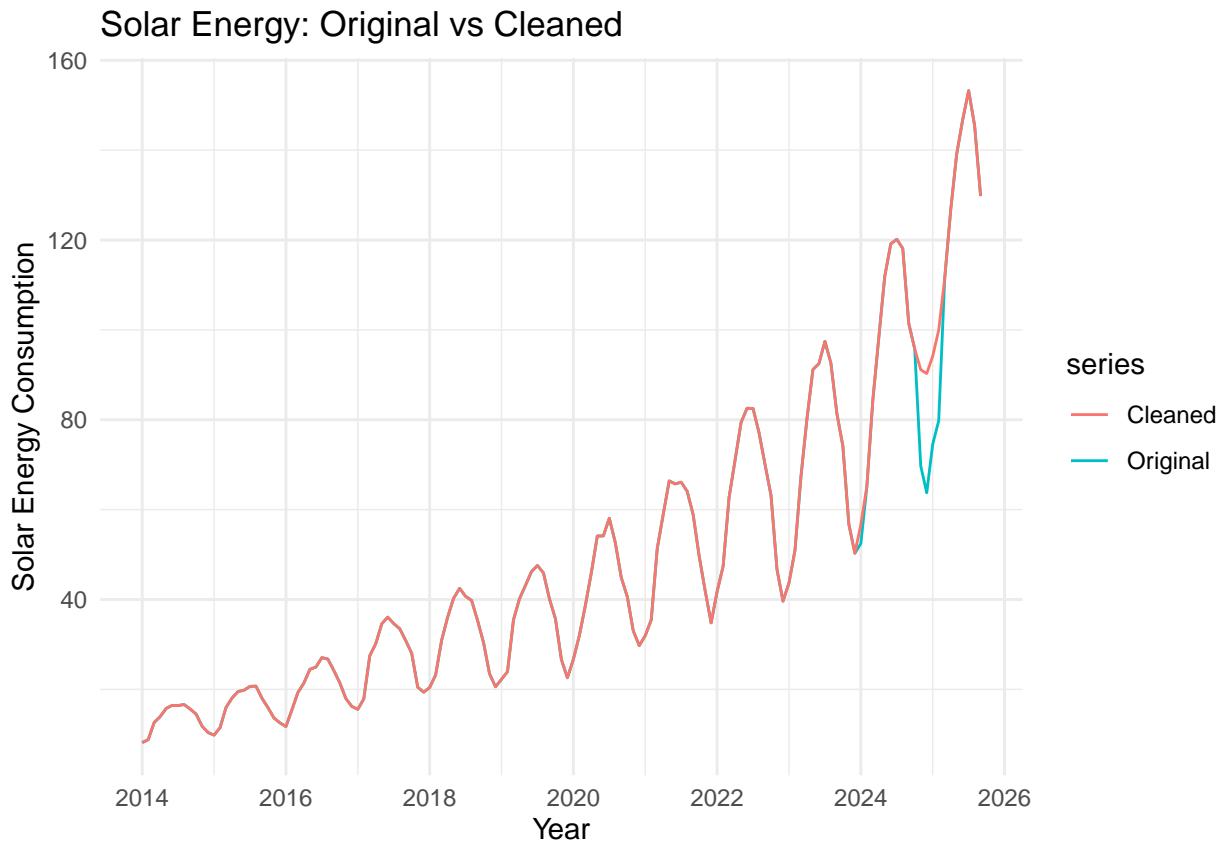
Identify and Remove outliers

Q8

Apply the `tsclean()` to both time series object you created on Q4. Did the function removed any outliers from the series? Hint: Use `autoplot()` to check if there is difference between cleaned series and original series.

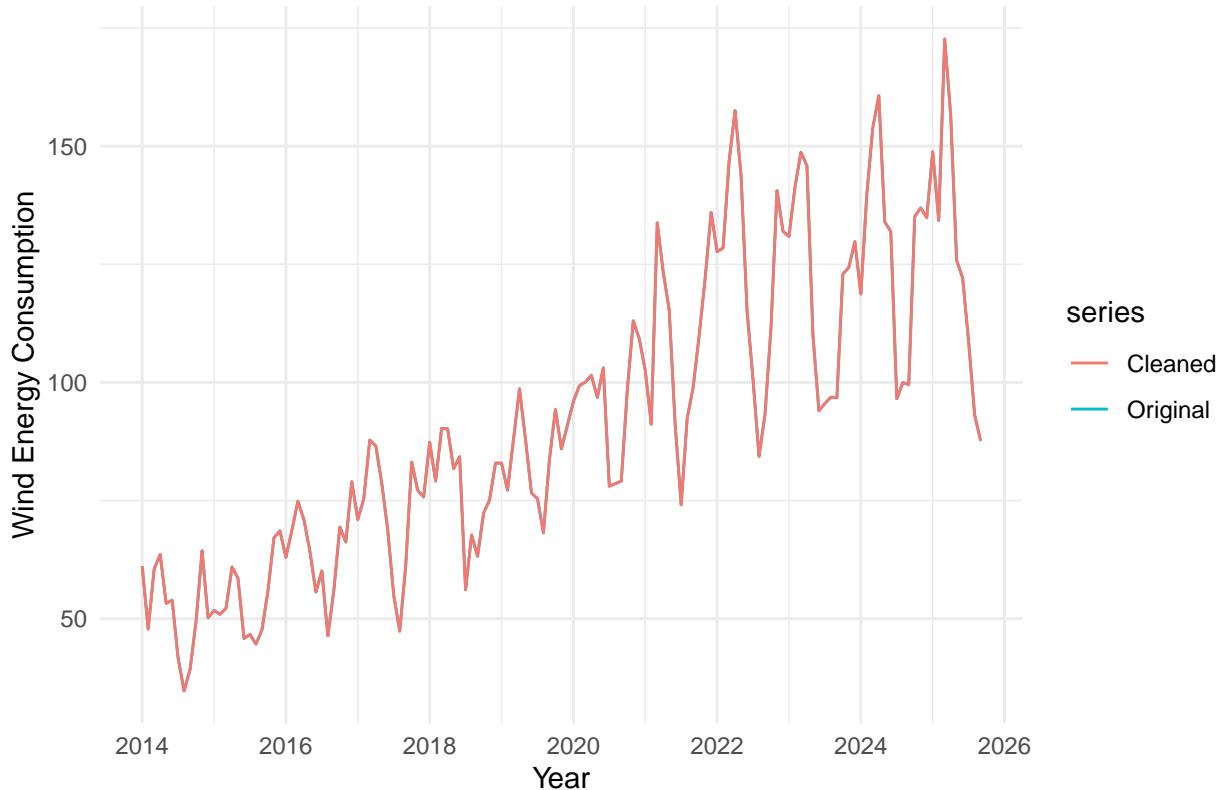
```
solar_recent_clean <- tsClean(solar_recent_ts)
wind_recent_clean <- tsClean(wind_recent_ts)

autoplot(solar_recent_ts, series = "Original") +
  autolayer(solar_recent_clean, series = "Cleaned") +
  xlab("Year") +
  ylab("Solar Energy Consumption") +
  ggtitle("Solar Energy: Original vs Cleaned") +
  theme_minimal()
```



```
autoplot(wind_recent_ts, series = "Original") +
  autolayer(wind_recent_clean, series = "Cleaned") +
  xlab("Year") +
  ylab("Wind Energy Consumption") +
  ggtitle("Wind Energy: Original vs Cleaned") +
  theme_minimal()
```

Wind Energy: Original vs Cleaned



Answer: There were outliers removed from both series, mainly after 2015 for solar energy consumption and after 2010 for wind energy consumption, which drastically decreased the variability of both for those timespans.

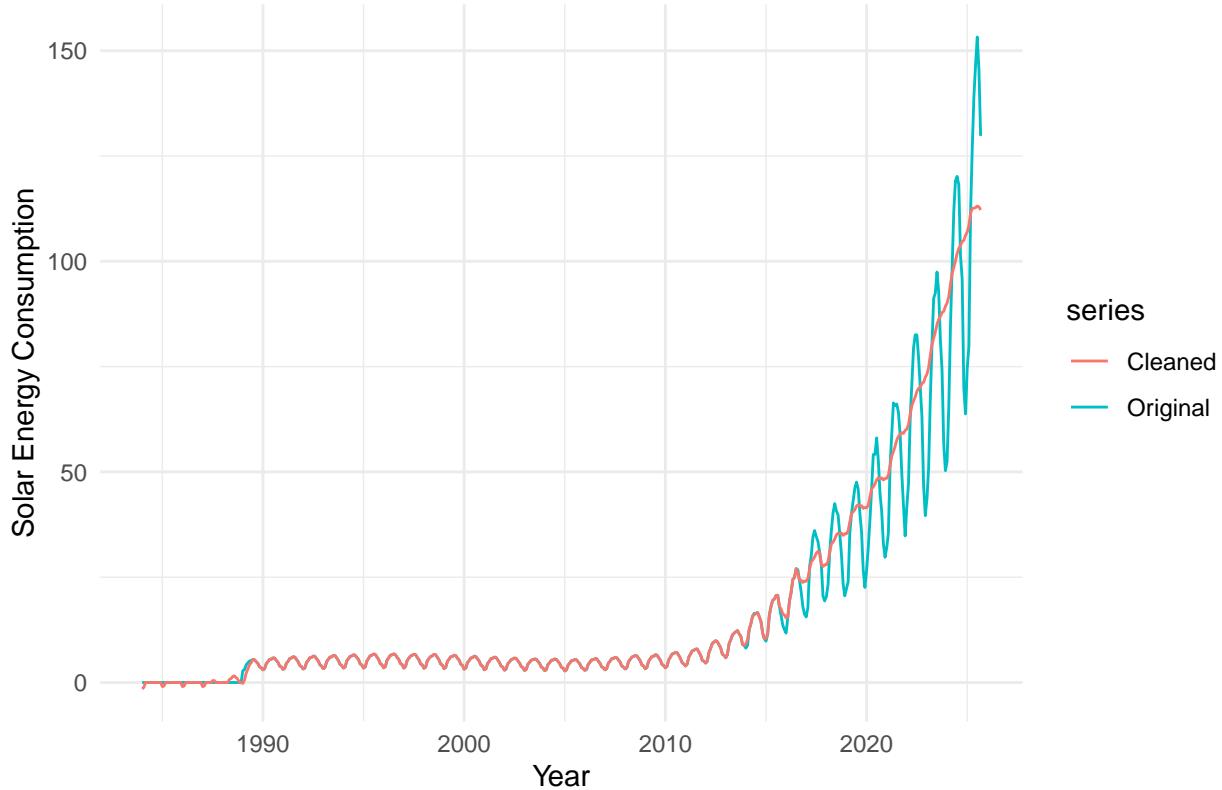
Q9

Redo number Q8 but now with the time series you created on Q7, i.e., the series starting in 2014. Using what `autoplot()` again what happened now? Did the function removed any outliers from the series?

```
solar_clean <- tsclean(solar_ts)
wind_clean <- tsapply(wind_ts)

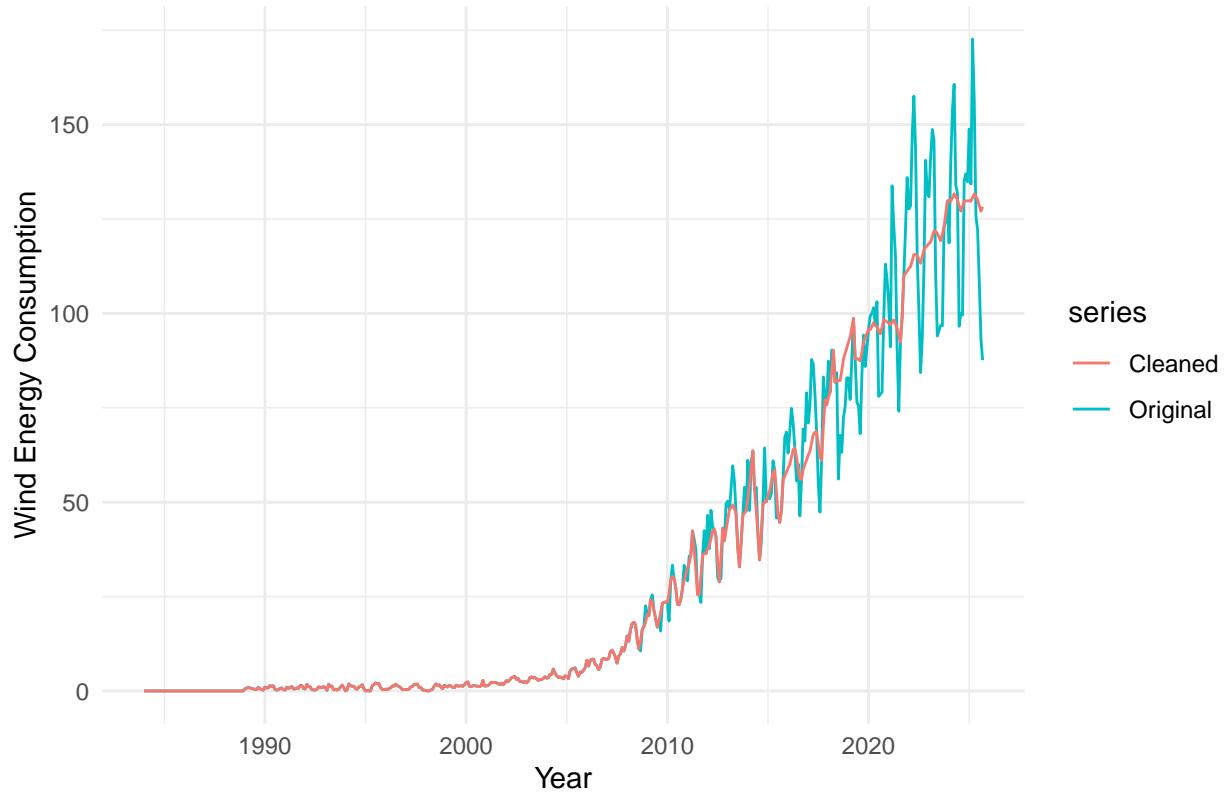
autoplot(solar_ts, series = "Original") +
  autolayer(solar_clean, series = "Cleaned") +
  xlab("Year") +
  ylab("Solar Energy Consumption") +
  ggtitle("Solar Energy: Original vs Cleaned") +
  theme_minimal()
```

Solar Energy: Original vs Cleaned



```
autoplot(wind_ts, series = "Original") +  
  autolayer(wind_clean, series = "Cleaned") +  
  xlab("Year") +  
  ylab("Wind Energy Consumption") +  
  ggtitle("Wind Energy: Original vs Cleaned") +  
  theme_minimal()
```

Wind Energy: Original vs Cleaned



Answer: For the more recent years, it seems that outliers were only removed from wind energy consumption, which did not drastically change the graph except for one year.